

Container + DeepLearning: from working to scaling

[Frank Zhao](#)

Software Architect, EMC CTO Office

zhaojp@gmail.com

WHO AM I?

- Frank Zhao, 赵军平, Junping Zhao
- Software Architect @ EMC CTO Office
- 10+year engineering in storage, virtualization, flash
 - 30+ patents (3 granted in U.S.)
- Now, working areas:
 - In-mem processing and analytics, perf acceleration
 - Micro-service, container
 - Streaming processing system
 - Software defined storage ...

AGENDA

- Motivations
- Deep learning and TensorFlow
- Distributed DL training
 - Docker + K8S + Tensorflow as example
- Distributed DL serving
 - Docker + K8S + Tensorflow as example
- Summary
- Outlook

Motivations

- Container + DL for
 - Easier deployment & management
 - Scalability: from infrastructure to typical DL framework
 - Smooth transition from training to serving, or Dev&Ops
 - Least performance overhead (vs. VM) especially CPU/mem
 - Maximize heterogeneous env, computing, memory, networking etc
- → **“DL as a Service” by container ecosystem?**
 - In cloud env, scale-out & heterogeneous
- Focus on DL workload, practices sharing, opportunities
 - Non-goal: TF internal or deep dive

Deep learning

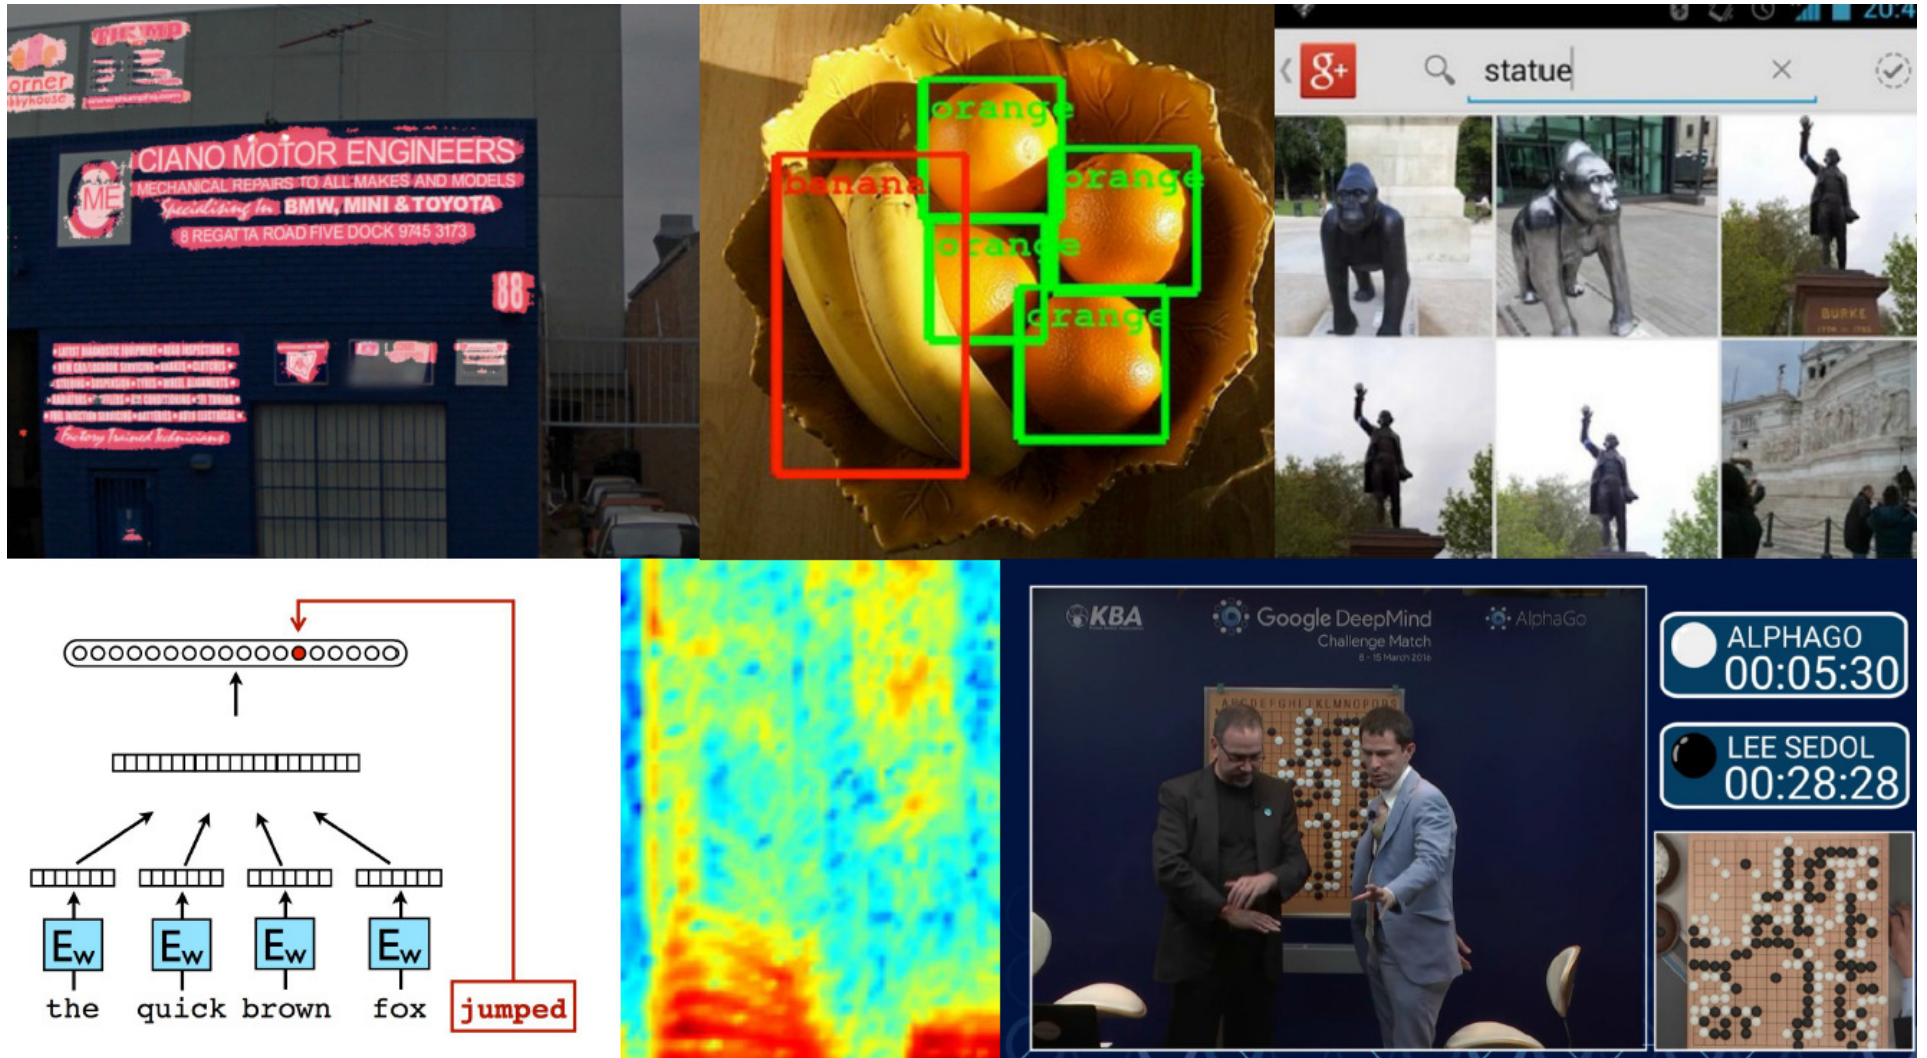


Photo from Google intro: “Diving into ML through Tensorflow”

Open-source DL frameworks


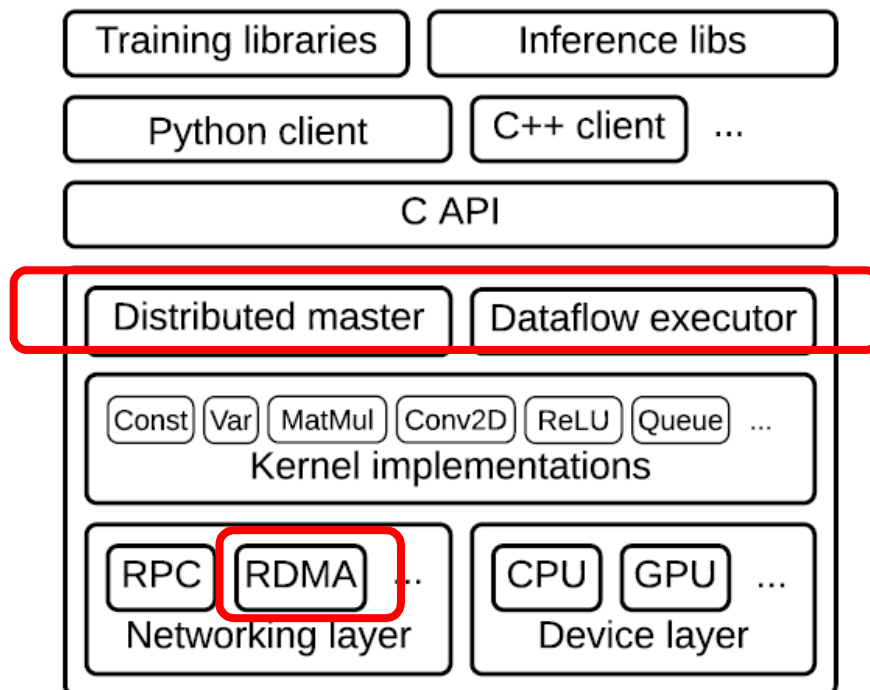
By 6/28/2016

	Star	Forks	Contributor
Tensorflow	26995	10723	286
Caffe	10973	6575	196
CNTK	5699	1173	69
Torch	4852	1360	100
Theano	4022	1448	234
MXNet	4173	1515	152
Apache SINGA	607	211	18

A more detailed summary and comparison @ <https://github.com/zer0n/deepframeworks>

TensorFlow

- DL framework from Google
 - GPU/CPU/TPU, heterogeneous: desktop, server, mobile etc
 - C++, Python; Distributed training and serving
 - DNN building blocks, ckpt/queue/ ..., TensorBoard
 - Good docker/K8S supported



TF distribution overview

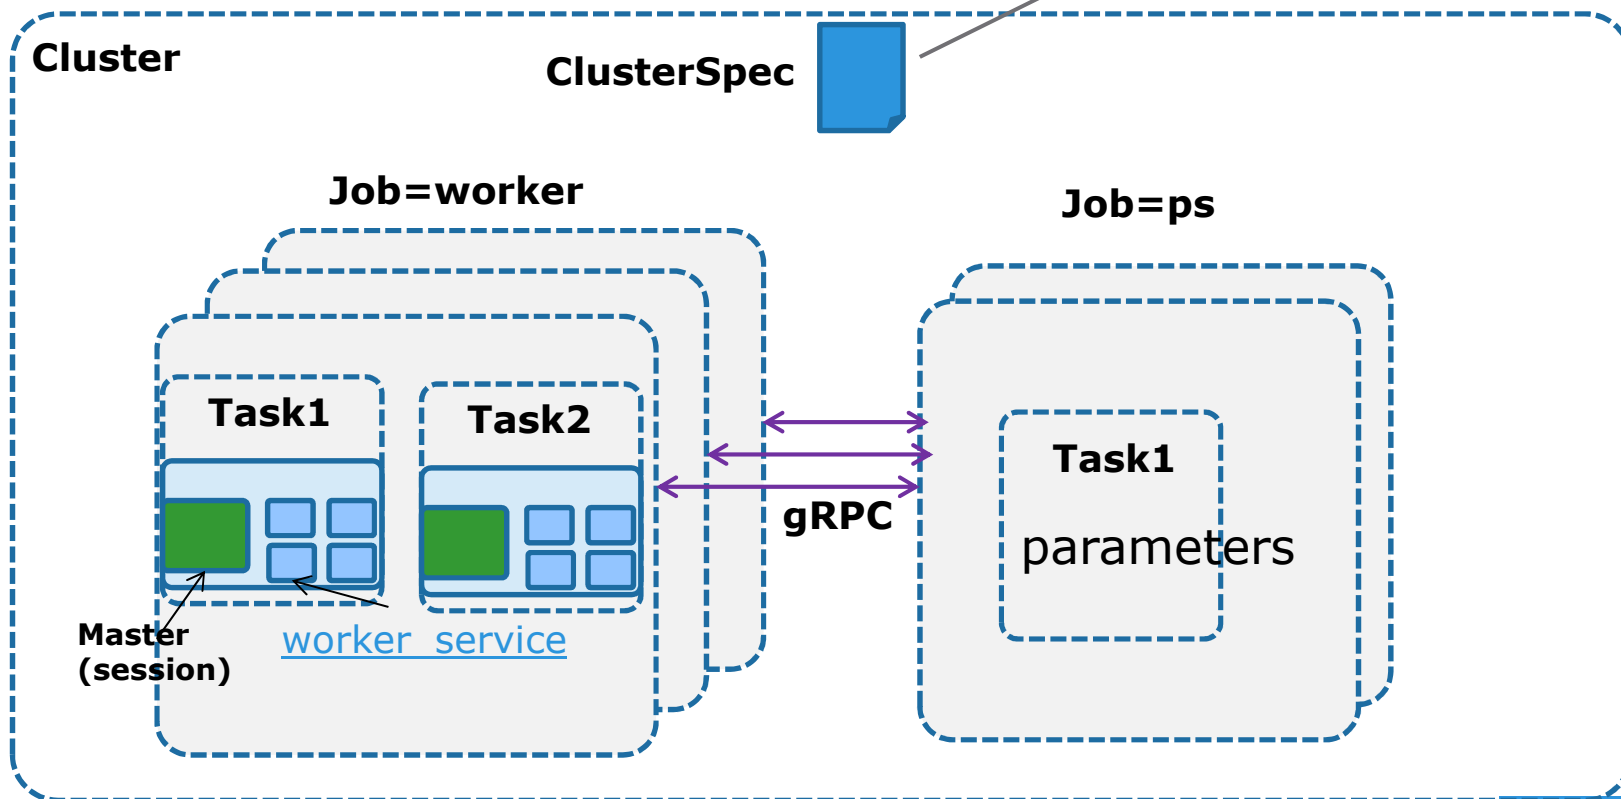
- **Cluster**

- **Jobs (ps, worker)**

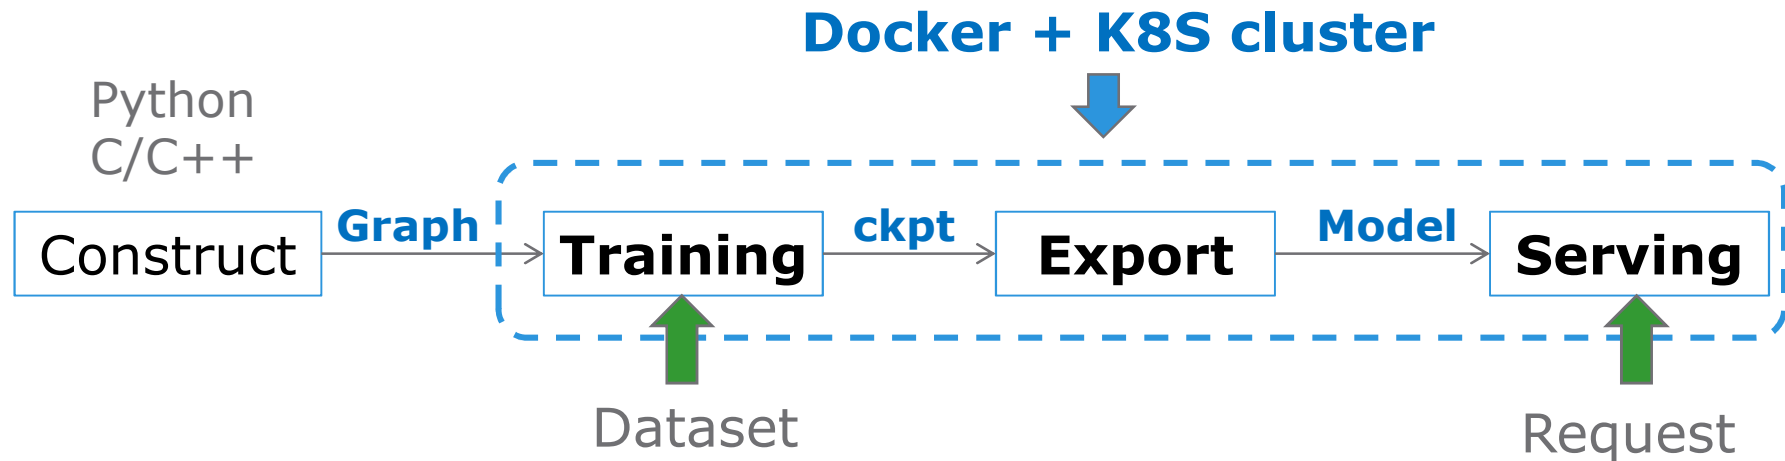
- **Tasks**

- **Devices** (CPU, GPU)

```
tf.train.ClusterSpec({  
  "worker": [  
    "worker0.example.com:2222",  
    "worker1.example.com:2222",  
    "worker2.example.com:2222"  
  ],  
  "ps": [  
    "ps0.example.com:2222",  
    "ps1.example.com:2222"  
  ]  
})
```



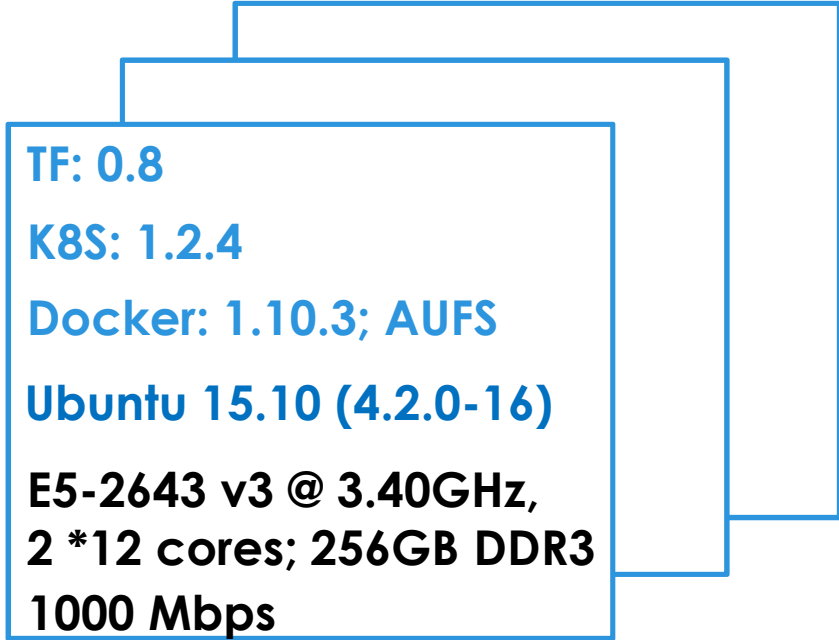
Workflow & dataset



Dataset	Type	Size
MNIST	Handwritten digit recognition	10 class. 60K training + 10K test
CIFAR-10/100	Image classification	10/100 classes; (50K + 10K test) 160+MB
Inception-V3	Image classification	1000 classes; ~500GB dataset

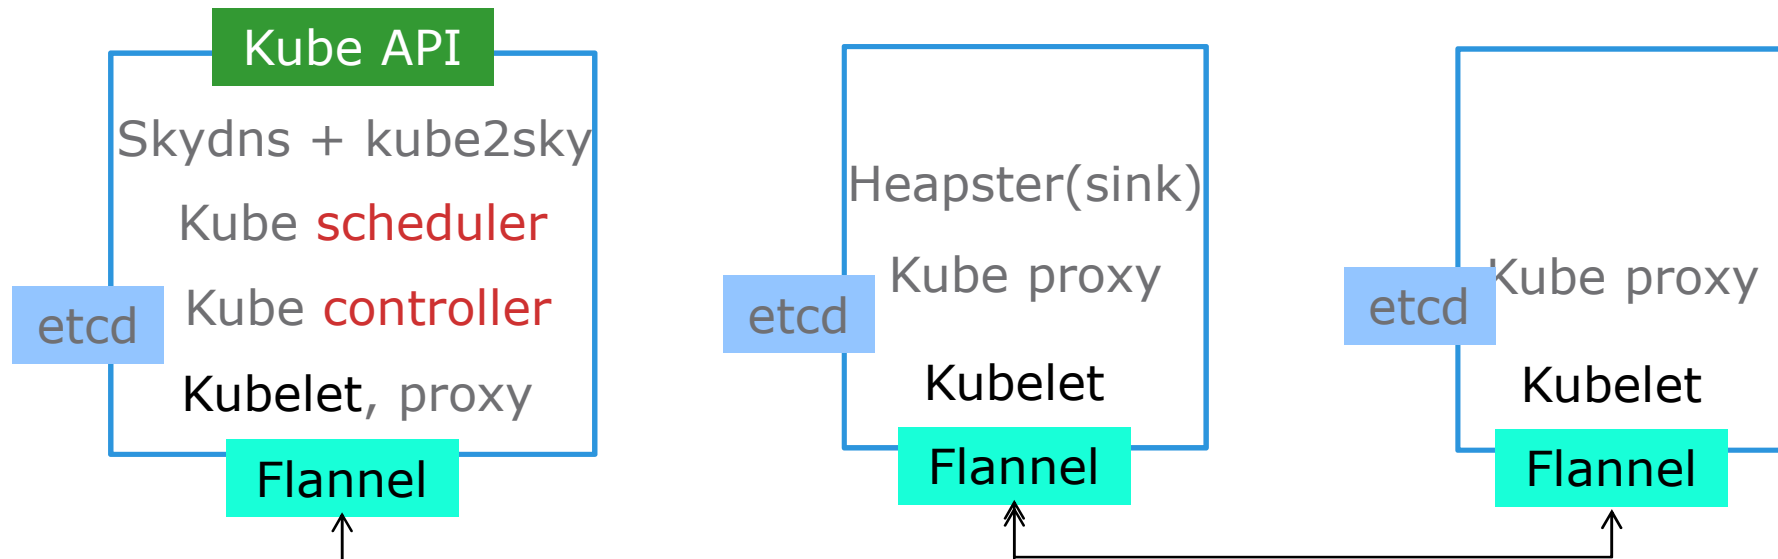
Test environment

- 3-node bare metal (Dell?) servers
 - CPU only; GPU+CPU cluster is being setup
 - Docker images from Tensorflow official w/ modifications



TF: 0.8
K8S: 1.2.4
Docker: 1.10.3; AUFS
Ubuntu 15.10 (4.2.0-16)
E5-2643 v3 @ 3.40GHz,
2 *12 cores; 256GB DDR3
1000 Mbps

K8S cluster



```

root@sm-tower1:~# docker ps
CONTAINER ID   IMAGE
6605cec68471   gcr.io/google_containers/heapster_grafana:v2.6.0-2
05102147185b   gcr.io/google_containers/heapster_influxdb:v0.5
e42bffd27ce5   gcr.io/google_containers/pause:2.0
90b53bc1341e   gcr.io/google_containers/skydns:1.0
527d600b9259   gcr.io/google_containers/kube2sky:1.15
de3633ab3fb0   quay.io/coreos/flannel:0.5.5
ff9c894c47e0   gcr.io/google_containers/flannel-server-helper:0.1
8b20a0d4a34b   quay.io/smana/kubernetes-hyperkube:v1.2.4
373f4ad80b61   quay.io/smana/kubernetes-hyperkube:v1.2.4
c661aa61186b   quay.io/smana/kubernetes-hyperkube:v1.2.4
4a4b30c7ccfb   gcr.io/google_containers/pause:2.0
04ec408af28e   gcr.io/google_containers/pause:2.0
bc41b8f1f32b   gcr.io/google_containers/pause:2.0
93e8cf10610c   gcr.io/google_containers/pause:2.0

COMMAND
"/bin/sh -c /run.sh"
"influxd --config /et"
"/pause"
"/skydns"
"/kube2sky"
"/bin/sh -c '/opt/bin"
"/flannel_helper --ne"
"/hyperkube proxy --v"
"/hyperkube scheduler"
"/hyperkube controlle"
"/pause"
"/pause"
"/pause"
"/pause"
"/pause"

CREATED
3 days ago
3 days ago
3 days ago
3 days ago
3 days ago
4 days ago
4 days ago
4 days ago
4 days ago
4 days ago
4 days ago
4 days ago
4 days ago
4 days ago
4 days ago

```

Cluster setup

- Still not very easy!
 - TF official setup script doesn't work in my Ubuntu15
- Tips
 - DNS (skydns) service is highly recommended
 - Use Docker 1.10 if possible
 - Docker1.11 (RunC) may not work well
 - Kargo likely helps setup easier
 - <https://docs.kubespray.io/> including K8S, etcd, flanneld, etc



- Close look at DL workload
 - Workload profiling
 - Scaling w/ CPU/GPU
- Distribution efficiency?
 - Communication overhead
 - Synchronization the parameter overhead
- Test plan: containers in **1**-node vs. in **3**-node
 - To isolate networking overhead
 - Async vs. sync

DL workload profiling

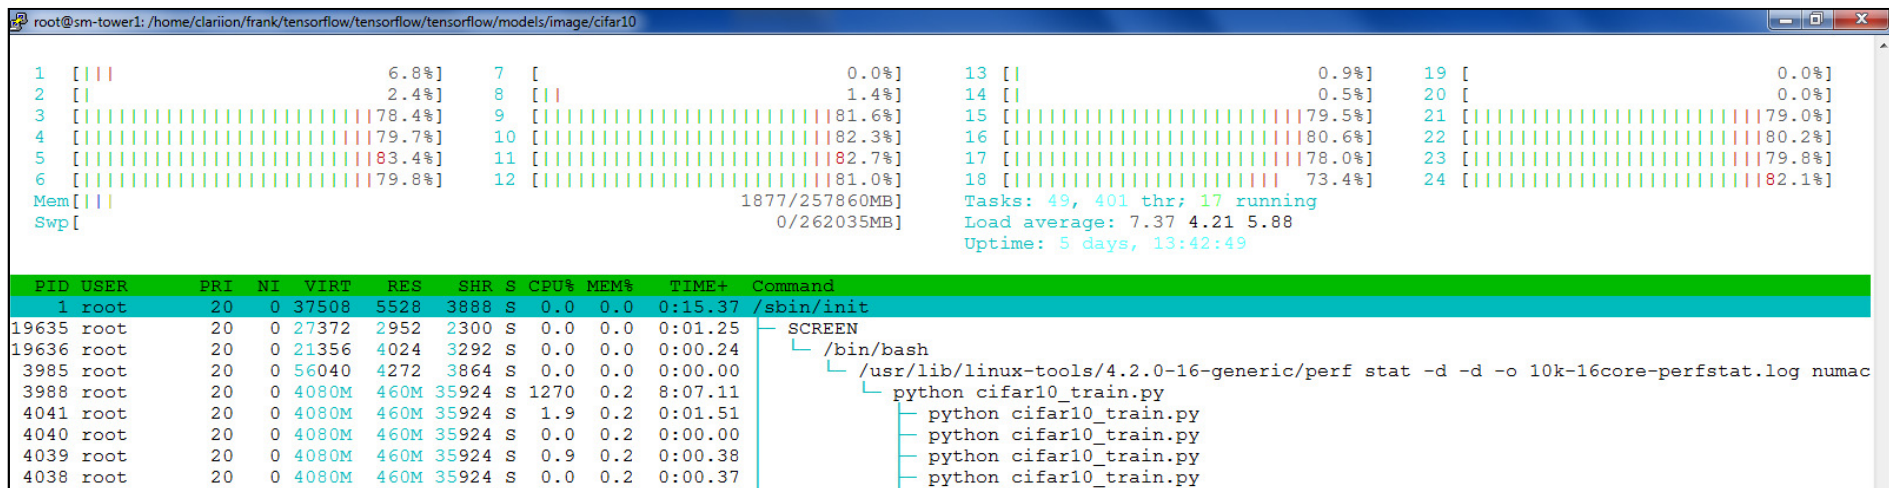
CIFAR10 dataset, CNN, Intel E5 v3 CPU*24 core

Performance counter stats for 'python cifar10_train.py':					
5263014.248826	task-clock (msec)	#	17.133 CPUs utilized		
22,264,167	context-switches	#	0.004 M/sec		
1,769,718	cpu-migrations	#	0.336 K/sec		
126,755,898	page-faults	#	0.024 M/sec		
18,645,619,419,538	cycles	#	3.543 GHz		(28.56%)
<not supported>	stalled-cycles-frontend				
<not supported>	stalled-cycles-backend				
21,654,529,625,343	instructions	#	1.16 insns per cycle		(36.24%)
1,170,696,486,325	branches	#	222.438 M/sec		(35.92%)
11,265,428,007	branch-misses	#	0.96% of all branches		(35.80%)
4,000,366,119,368	L1-dcache-loads	#	760.090 M/sec		(21.86%)
327,032,499,068	L1-dcache-load-misses	#	8.18% of all L1-dcache hits		(16.77%)
53,063,949,994	LLC-loads	#	10.082 M/sec		(16.70%)
11,964,635,507	LLC-load-misses	#	45.10% of all LL-cache hits		(21.86%)
<not supported>	L1-icache-loads				
12,139,313,035	L1-icache-load-misses	#	2.307 M/sec		(28.53%)
3,996,507,473,413	dTLB-loads	#	759.357 M/sec		(21.46%)
4,198,057,989	dTLB-load-misses	#	0.11% of all dTLB cache hits		(18.96%)
1,276,700,163	iTLB-loads	#	0.243 M/sec		(16.68%)
624,640,855	iTLB-load-misses	#	48.93% of all iTLB cache hits		(21.60%)

- **CPU intensive, consume 17+ cores**
- **Cache locality is not good**

Balancing btw CPUs


- Very good balancing btw cores



Performance Scaling w/ CPU

- 4X cores gains 1.8X faster
 - looks ok, but not quite linear. rooms to improve?

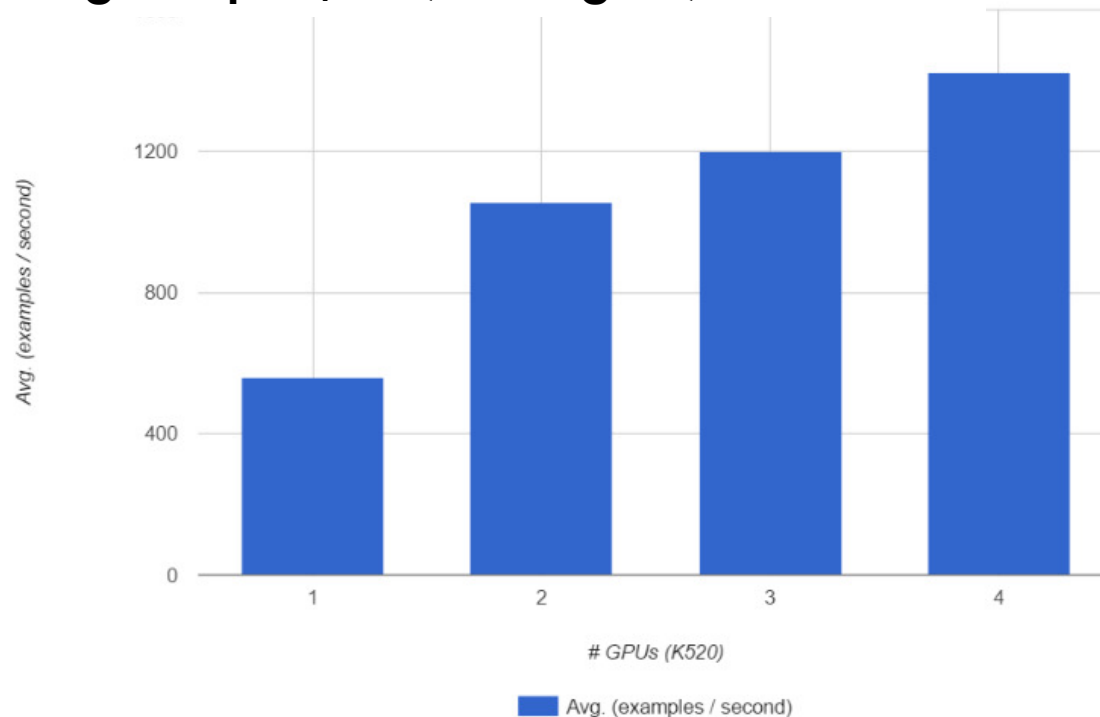
CNN-CIFAR10 training w/ multiple CPU cores



Performance Scaling w/ GPU

- 4X GPU achieves 2.4X faster
 - AWS g2.8xlarge, 4*K520 GPU

Avg samples/sec, the higher, the better



benchmark from <http://www.bitfusion.io/2016/05/09/easy-tensorflow-model-training-aws/>

Distribution: Parallelism & parameter update

Data parallel (Full-graph)	Model parallel (Btw-graph)
Deterministic model. Same codes in nodes/workers	fine-grained control of ops on node/workers; may different
Simple, easy to implement & control	Complex, may with high accuracy

Synchronous update	Asynchronous update
sync-point between workers. Prevents from "falling behind"	Worker runs as fast as they can w/o sync-point
high quality/accuracy, but slower* (or to optimize)	Usually fast, but relative lower accuracy due to fall-behind

Distributed training steps, by Docker+K8S

1. Construct the graph,
 - full-graph (data parallel) for Inception and MNIST
 - Sync or async; place worker on CPU or GPU
 - Setup termination condition (training steps or accuracy)
2. Package as container
 - Graph model, TF running lib, gRPC, or even dataset
3. Deploy by K8S
 - Specify job/worker and index
4. Feed dataset and start training
 - Container bind-volume, and batch load
5. Monitoring, visualization & report




Training cluster config

- K8S config yaml

- tools/dist_test/scripts/k8s_tensorflow.py

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: tf-worker0
spec:
  replicas: 1
  template:
    metadata:
      labels:
        tf-worker: "0"
    spec:
      containers:
        - name: tf-worker0
          image: tensorflow/tf_grpc_test_server
          args:
            - --cluster_spec=worker|tf-worker0:2222;tf-worker1:2222;tf-worker2:2222,
            - --job_name=worker
            - --task_id=0
          ports:
            - containerPort: 2222
      nodeSelector:
        nodeName: "tower1"
```




Docker image w/
model + TF+gRPC

Cluster spec;
specify PS|worker and index

Distribute containers in **1** node

- Scheduled 3 containers in 1 node using K8S "nodeSelector"
- Async vs. sync

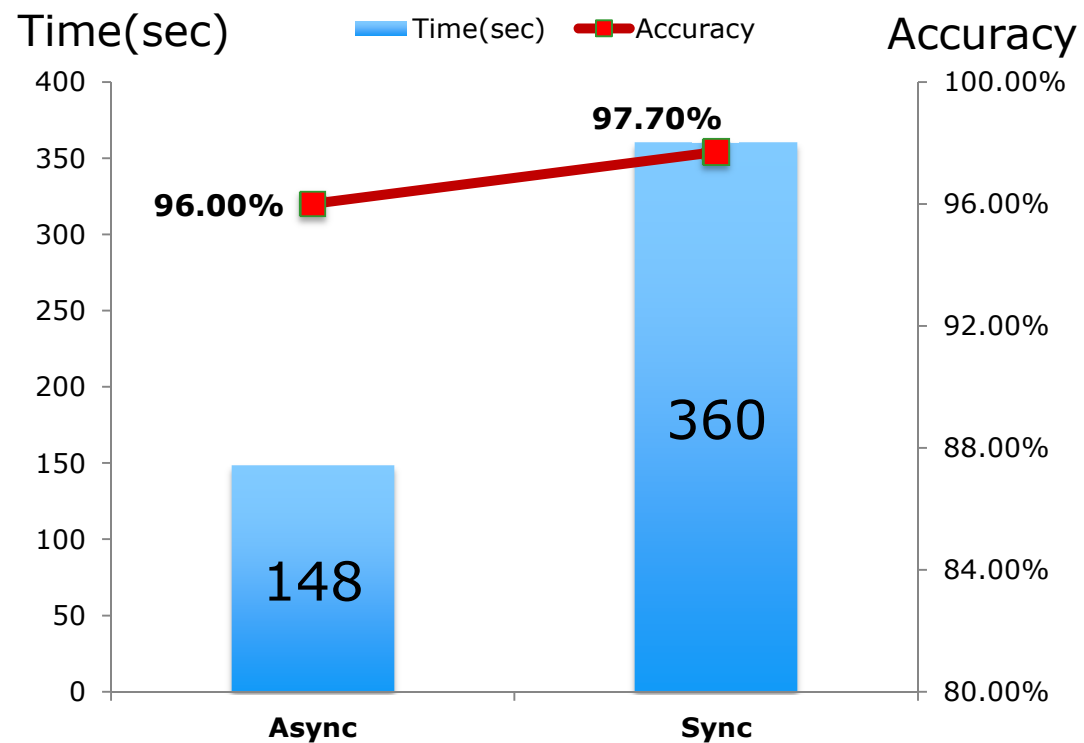


Distribute containers in 3 nodes


Each node has 1 worker container

- Async vs. sync

MNIST training time and accuracy



Running steps in workers



Observation and Considerations

- Async vs. Sync, Efficiency
 - Async is much faster than sync (**68%** in one node)
 - 143% in 3 nodes, but network is weak and stack is not fully optimized
 - Communication and sync overhead are bottlenecks
 - Fine tuned distributed training can achieve **13X faster w/ 16 nodes***
 - **RDMA**, parameter **local cache** + **bulk async update** back
 - Spark and Nvidia HPC etc claimed get 50+% higher perf
- Data parallel
 - Partition dataset across workers and proper mini-batch
 - shared bind-volume across workers
 - Not balanced in async way - to dig out
- Pod placement on nodes
 - Automatic, or use K8S nodeSelector/nodeAffinity
- Ops placement on CPU/GPU device
 - Automatic, or explicit TF API: with tf.device(name): {ops}

Observation and Considerations(2)

- Instance number
 - Moderate num per node to reduce communication overhead
- CPU/Mem throttling
 - Config “Request/Limit” only if necessary i.e., resource-shared env.
 - Container level, K8S pod, K8S replica or quota
 - Carefully set batch size to avoid OOM
 - Fully exploit GPU mem at first
- Networking concern, Flannel VLAN overhead
 - Use “--net=host” with native network and map to ports
 - RDMA is recommended; need powerful HW/driver
- With GPU?
 - Nvidia-Docker-plugin
 - GPU memory is fast but capacity is limited - a few GB
 - **Mem tiering:** GPU mem as Tier1, swap to CPU mem
 - Reference: “Scalable deep learning on distributed GPUs with a GPU-specialized parameter server” EuroSys16

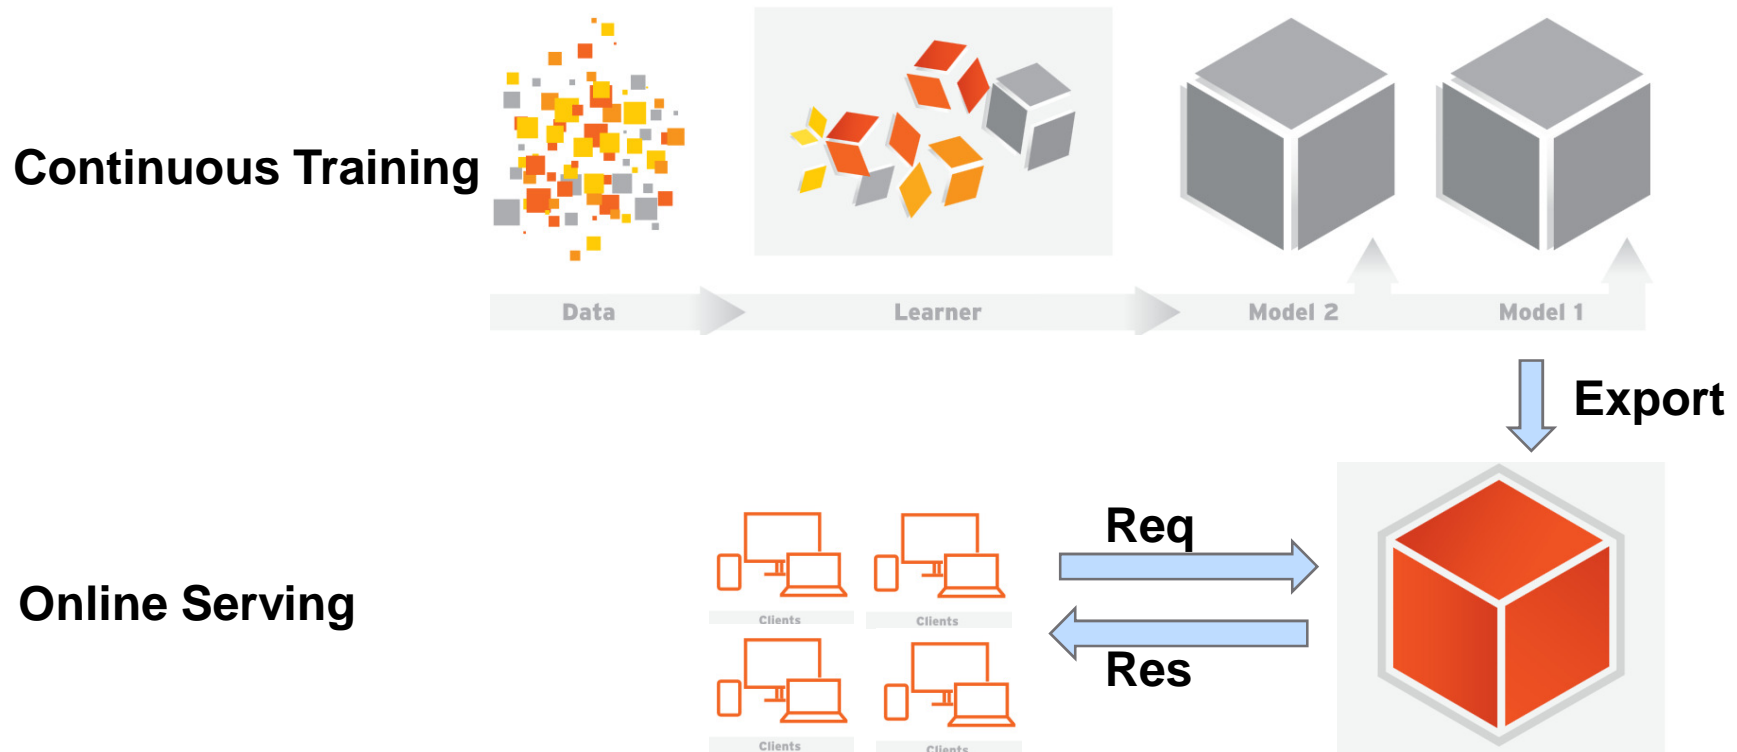
Export model

- Checkpoint then export it
 - saver() : create a ckpt into files
 - Exporter(): set signature, version etc
 - CreateSessionBundle(): to load model at specific path

```
root@:/serving# ll inception-export
total 107M
-rw-r--r--    69 Jun 21 09:25 checkpoint
-rw-r--r-- 107M Jun 21 09:25 export
-rw-r--r-- 140K Jun 21 09:25 export.meta
```

DL serving

- Online service to the world based on trained model
 - No label dataset, no loss function and backpropagation



Distributed Serving by Docker + K8S

- K8S features for deployment and management
 - Deployment
 - Auto-scale
 - Rolling-upgrade

Prepare the image

- + Include tensorflow,
- + tensorflow_serving
- + exported model/parameter

Export service port

Example: `gcr.io/tensorflow-serving/inception`

K8S: "deployment"

- **Highly recommended**, rather than k8s RC, pod
 - Auto manage desired state, i.e, autoscale (HPA)
 - With service and Replicaset
 - Easy to rolling upgrade on the fly

```
{
  "apiVersion": "extensions/v1beta1",
  "kind": "Deployment",
  "metadata": {
    "name": "inception.com",
    "namespace": "kube-system"
  },
  "spec": {
    "replicas": 1,
    "template": {
      "spec": {
        "containers": [
```

Same ns w/ heapster

```
    "name": "inception-container",
    "image": "inception_serving:v1",
    "imagePullPolicy": "IfNotPresent",
    "command": [
      ],
    "resources": {
      "requests": {
        "cpu": "6000m"
      }
    }
  }
```

Required for auto-scale

K8S: auto-scale

- Need Heapster (Influxdb) to collect matrix

```
# kubectl get pods --all-namespaces -o wide
```

NAME	READY	STATUS	AGE	NODE
heapster-v1.0.2-3098241085-1yv64	4/4	Running	1h	sm-tower2
monitoring-influxdb-grafana-v3-pjib8	2/2	Running	2h	sm-tower1
inception.com-2008606290-16akw	1/1	Running	29s	sm-tower1
inception.com-2008606290-aluae	1/1	Running	29s	sm-tower2
inception.com-2008606290-va0co	1/1	Running	29s	sm-tower0
kube-controller-manager-sm-tower0	1/1	Running	14d	sm-tower0

kubectl autoscale deployment inception.com

--min=1 --max=3 --cpu-percent=50 --namespace="kube-system"

K8S: auto-scale result

- Increase workload, CPU get hot

```
# kubectl get hpa --namespace="kube-system"
```

NAME	REFERENCE	TARGET	CURRENT
inception.com	Deployment/inception.com/scale	50%	878%

Auto-scale# ← in 4+ minutes

kubectl describe deployments inception.com

LastSeen	Count	Reason	Message
-----	-----	-----	-----
4m	1	ScalingReplicaSet	Scaled up replica set inception.co

K8S: rolling upgrade

- Minimize service impact via online rolling upgrade
- "apply"
 - Replace instance in graceful pace

```
"spec": {  
  "containers": [  
    {  
      "name": "inception-container",  
      "image": "inception_serving:v2",  
      "imagePullPolicy": "IfNotPresent",  
      "ports": [  
        {  
          "containerPort": 80,  
          "protocol": "TCP"  
        }  
      ]  
    }  
  ]  
}
```

kubectl **apply** -f inception-v2.json --validate=false
deployment "inception.com" configured
service "inception-service" configured

K8S: rolling-upgrade in-progress

- `kubectl get pods --namespace="kube-system" -o wide`

NAME	STATUS	RESTARTS	AGE	NODE
<code>inception.com-2008606290-0w8w2</code>	Terminating	0	8m	<code>sm-tower1</code>
<code>inception.com-2008606290-2kxou</code>	Terminating	0	37m	<code>sm-tower0</code>
<code>inception.com-2008606290-vq977</code>	Terminating	0	8m	<code>sm-tower2</code>
<code>inception.com-2103174739-1cjic</code>	Running	0	18s	<code>sm-tower0</code>
...				
<code>inception.com-2103174739-1cjic</code>	Running	0	14m	<code>sm-tower0</code>
<code>inception.com-2103174739-4kws4</code>	Running	0	34s	<code>sm-tower1</code>
<code>inception.com-2103174739-ypwlc</code>	Running	0	4m	<code>sm-tower2</code>

Considerations, and gaps

- Overall, K8S deployment is wonderful for serving
- Pod placement, re-scheduling
 - Use nodeSelector or nodeAffinity, label based
 - Gap: only for initial placement, not for dynamic execution
- Auto-scaling matrix and rule
 - Now CPU usage based, need heapster config
 - Gap: need GPU support and advanced rules
- Auto scale-down?
 - K8S supports, but not stable in my test.
 - Use “resize” cmd to manually adjust
 - What if scaling meets auto-upgrade??

Summary

- Explore typical DL workload
- Practice end-end DL distributed training and serving
 - Docker + K8S + TF
 - More deep optimization work are in-progress, see next
- TF, as one of hottest DL frameworks, provides good docker/K8S support, distributed version is promising but still in early stage
 - Perf scaling w/ multiple CPU/GPU can be improved
 - Data distribution is not always balanced
 - Lack of deep guideline, and profiling/tuning module (“EEG”) is not open-sourced
- K8S: auto-scale based on GPU usage or mixed?
-

In-progress & outlook

- Community

- TF and Google

- C++/Python optimization; More platforms (iOS, Windows, OpenCL etc)
 - Optimization for distribution, memcpy btw host and GPU etc
 - Google services: Datalab, dataflow(Apache Beam), cloud ML

- K8S: advanced matrix for auto-scale; node affinity

- Our in-progress/plan

- Heterogeneous cluster w/ powerful HW

- GPU + CPU, 100GbE+RDMA/GPUDirect, NVRAM/NVMe

- Unified memory and tiering for parameter update

- Smart container/ops placement and (re-)scheduling

- Efficient data partition/paralleling, like HPC

→ **“DL as a Service”**

References

- TensorFlow: A system for large-scale machine learning; on Heterogeneous Distributed Systems
 - <http://arxiv.org/abs/1605.08695> and <http://bit.ly/tf-workshop-slides>
- K8S and TF serving
 - <http://blog.kubernetes.io/2016/03/scaling-neural-network-image-classification-using-Kubernetes-with-TensorFlow-Serving.html>
- [GeePS: Scalable deep learning on distributed GPUs with a GPU-specialized parameter server](#) EuroSys 2016
- Google cloud DL services, TPU
 - <http://conferences.oreilly.com/strata/hadoop-big-data-ca/public/schedule/detail/50445>
 - <https://cloudplatform.googleblog.com/2016/05/Google-supercharges-machine-learning-tasks-with-custom-chip.html>
- Scaling TF on Spark
 - <https://spark-summit.org/east-2016/events/distributed-tensor-flow-on-spark-scaling-googles-deep-learning-library/>
- Docker SWARM vs K8s
 - <https://blog.docker.com/2016/03/swarmweek-docker-swarm-exceeds-kubernetes-scale/>

Thank You! 谢谢!

ありがとう!

Questions?